

Lecture 04  
2020 Spring Data-622  
Bias Variance Tradeoff  
Raman Kannan

**Instructor Email Address:** [Raman.Kannan@sps.cuny.edu](mailto:Raman.Kannan@sps.cuny.edu)


Acknowledgements:

Generous support from IBM Power Systems Academic Initiative  
IBM PSAI provides computing infrastructure for free

# Bias/Variance

Error, because we don't have representative sample to learn from. Need more data. Estimating required data size is an important exercise in statistics.

Figure 1. Bias and variance in dart-throwing.



Error because the model is too simple, underfitting, Model is so simple that it cannot fit even when labels are given, training data.

Manifests as Error during training is significant  
Increase model complexity.

Error because the model captures even noise from the training data and is unable to determine class with never seen before data. Very small error during Training phase and much larger error during Testing phase (never seen before data).

Note error(training) is always lesser than error(testing),

<https://homes.cs.washington.edu/~pedrod/papers/cacm12.pdf>

# Variance Reduction Strategy

## Cross Validation

Instead of 80% for training and 20% for testing,

Create equally sized subsets of data,  
Iterate (train,test) over all the subsets, keeping one subset as test data  
In each iteration.

Use the average of the model. (model parameter average)

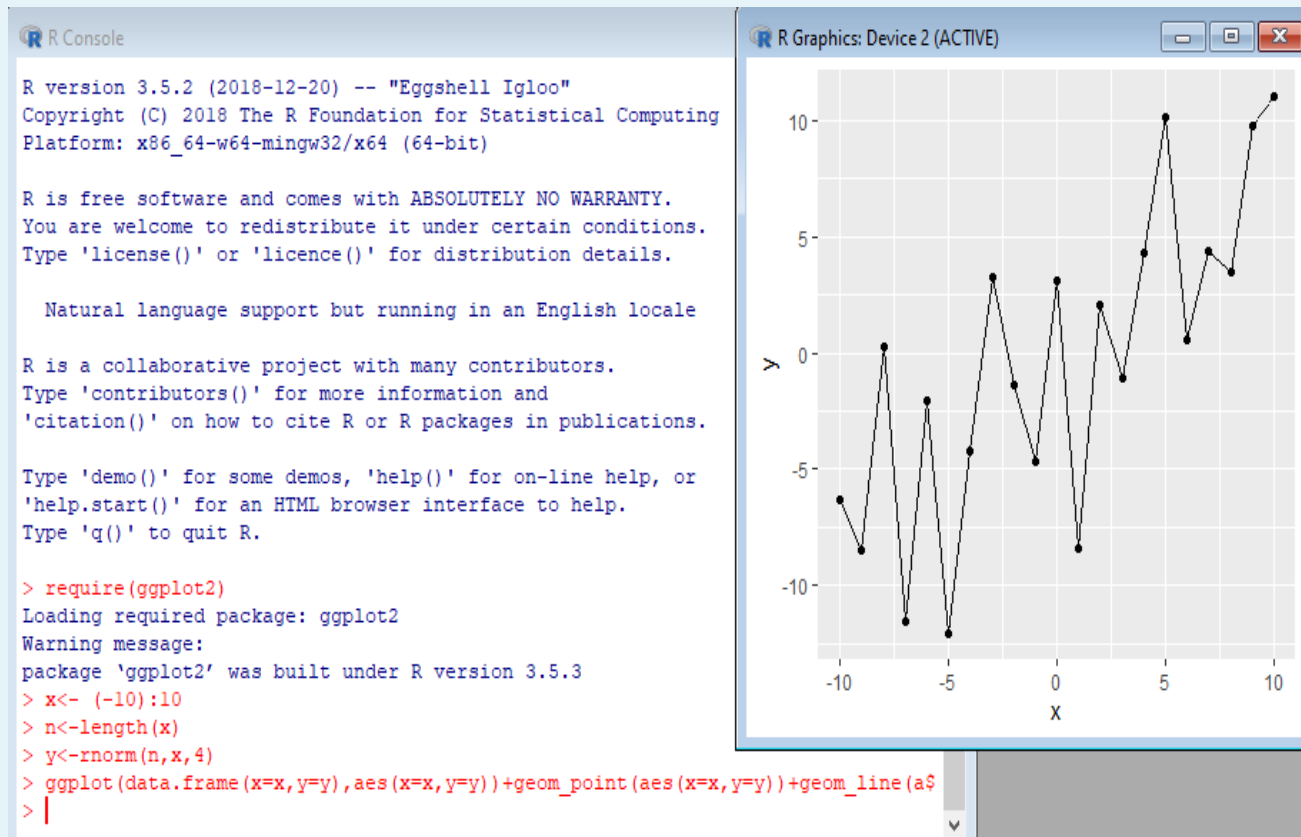
10-fold CV or 5-fold CV or 3-fold CV.

Or Leave one out (LOO) CV.

LOO CV: For each observation, exclude that observation, train on the  
Rest of the data, test on the excluded observation  
Finally take the average. Compute intensive.

Note with any type of CV, every available observation is used to train and  
Test unlike the traditional process – where some% of data is excluded  
from training and training.

## Overfitting – high variance

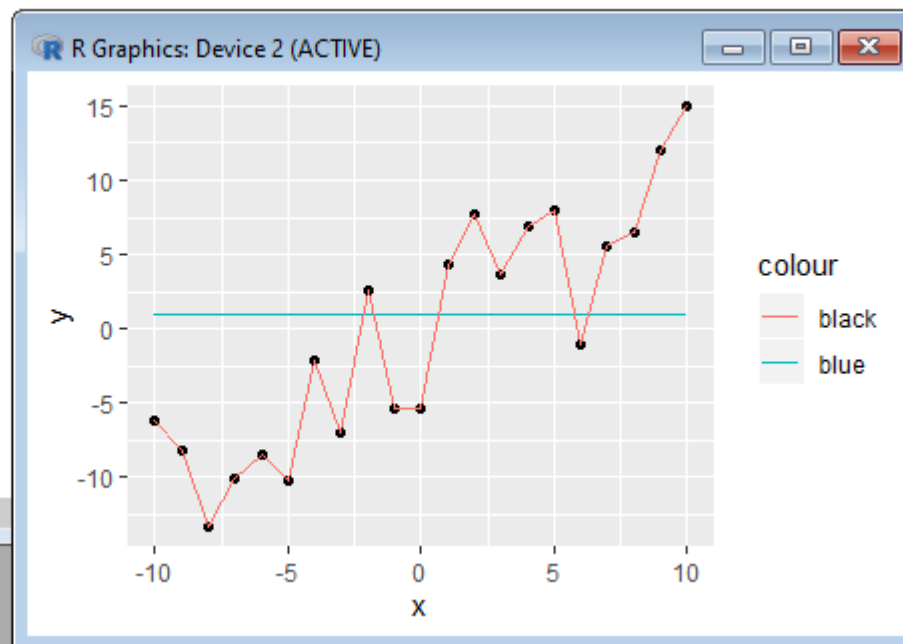


```
x<- (-10):10
n<-length(x)
y<-rnorm(n,x,4)
ggplot(data.frame(x=x,y=y),aes(x=x,y=y))+geom_point(aes(x=x,y=y))+geom_line(aes(x=x,y=y))
```

## Underfitting – high bias

```
x<- (-10):10  
n<-length(x)  
y<-rnorm(n,x,4)  
ggplot(data.frame(x=x,y=y),aes(x=x,y=y))+geom_line(aes(x=x,y=1))+geom_point(aes(x=x,y=y))+  
geom_line(aes(x=x,y=y))
```

```
> ggplot(data.frame(x=x,y=y), aes(x=x,y=y)) + geom_line(aes(x=x,y=1, color="blue")) + geom_point(aes(x=x,y=y)) +  
+ geom_line(aes(x=x,y=y, color="black"))  
> |
```




# Smooth – less variability and lower bias

```
x<- (-10):10  
n<-length(x)  
y<-rnorm(n,x,4)
```


```
ggplot(data.frame(x=x,y=y),aes(x=x,y=y))+geom_point(aes(x=x,y=y))+geom_smooth()+  
geom_line(aes(x=x,y=y))
```

```
> ggplot(data.frame(x=x,y=y),aes(x=x,y=y))+geom_point(aes(x=x,y=y))+geom_smooth()+ geom_line(aes(x=x,y=y))  
`geom_smooth()` using method = 'loess' and formula 'y ~ x'  
> |
```



# Shrinkage (aka Penalization)

<https://towardsdatascience.com/bias-variance-and-regularization-in-linear-regression-lasso-ridge-and-elastic-net-8bf81991d0c5>



# L-1 Lasso

<https://towardsdatascience.com/bias-variance-and-regularization-in-linear-regression-lasso-ridge-and-elastic-net-8bf81991d0c5>

<https://towardsdatascience.com/bias-variance-and-regularization-in-linear-regression-lasso-ridge->  
NY EMAIL ISL M Co Statistics How to Stop Skype Fr Co Attributes: Defining Cl Fundry & Yesco - Firm Mo

## Lasso

$$\min_{\beta \in \mathbb{R}^p} \left\{ \frac{1}{N} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \right\}$$

The Lasso cost function, from Wikipedia

Lasso (sometimes stylized as LASSO or lasso) adds an additional term to the cost function, adding the sum of the coefficient values (the L-1 norm) multiplied by a constant

lambda. This additional term penalizes the model for having coefficients that do not explain a sufficient amount of variance in the data. It also has a tendency to set the coefficients of the bad predictors mentioned above 0. This makes Lasso useful in feature selection.

Lasso however struggles with some types of data. If the number of predictors (p) is greater than the number of observations (n), Lasso will pick at most n predictors as non-zero, even if all predictors are relevant. Lasso will also struggle with colinear features (they're related/correlated strongly), in which it will select only one predictor to represent the full suite of correlated predictors. This selection will also be done in a random way, which is bad for reproducibility and interpretation.

It is important to note that if lambda=0, we effectively have no regularization and we will get the OLS solution. As lambda tends to infinity, the coefficients will tend towards 0 and the model will be just a constant function.

Feature selection



# L-2 Ridge

<https://towardsdatascience.com/bias-variance-and-regularization-in-linear-regression-lasso-ridge-and-elastic-net-8bf81991d0c>

## Ridge Regression

$$\hat{\beta}^{\text{ridge}} = \underset{\beta \in \mathbb{R}}{\operatorname{argmin}} \|y - XB\|_2^2 + \lambda \|B\|_2^2$$

Thanks to Kyoosik Kim

Ridge regression also adds an additional term to the cost function, but instead sums the squares of coefficient values (the L-2 norm) and multiplies it

by some constant lambda. Compared to Lasso, this regularization term will decrease the values of coefficients, but is unable to force a coefficient to exactly 0. This makes ridge regression's use limited with regards to feature selection. However, when  $p > n$ , it is capable of selecting more than  $n$  relevant predictors if necessary unlike Lasso. It will also select groups of colinear features, which its inventors dubbed the 'grouping effect.'

Much like with Lasso, we can vary lambda to get models with different levels of regularization with lambda=0 corresponding to OLS and lambda approaching infinity corresponding to a constant function.

Interestingly, analysis of both Lasso and Ridge regression has shown that neither technique is consistently better than the other; one must try both methods to determine which to use (Hou, Hastie, 2005).

# Combining L1 and L2 norms

## Elastic Net

$$\hat{\beta} \equiv \underset{\beta}{\operatorname{argmin}} (\|y - X\beta\|^2 + \lambda_2 \|\beta\|^2 + \lambda_1 \|\beta\|_1).$$

Thanks to Wikipedia

Elastic Net includes both L-1 and L-2 norm regularization terms. This gives us the benefits of both Lasso and Ridge regression. It has been found to

have predictive power better than Lasso, while still performing feature selection. We therefore get the best of both worlds, performing feature selection of Lasso with the feature-group selection of Ridge.

Elastic Net comes with the additional overhead of determining the two lambda values for optimal solutions.

<https://medium.com/@jayeshbahire/lasso-ridge-and-elastic-net-regularization-4807897cb722>

# OLS Basic Regression

Load data  
Split Training  
Split Testing  
Create Model  
Find Training Error  
Predict Testdata  
Estimate testing error  
Run Regularization  
Compare

```
rsquare <- function(given, predicted) {  
  sse <- sum((predicted - given)^2)  
  sst <- sum(given^2)  
  rsq <- 1 - sse / sst  
  
  # For this post, impose floor...  
  if (rsq < 0) rsq <- 0  
  
  return (rsq)  
}  
msd<-function(given,predicted) {  
  sqrt(mean((given-predicted)^2))  
}  
mtcarshp<-mtcars[,c("mpg","wt","drat","hp")]  
  
ntrain<-round(0.8*nrow(mtcars))  
ntest<-nrow(mtcars)-ntrain  
allidx<-1:nrow(mtcars)  
trainidx<-sample(allidx,ntrain,rep=FALSE)  
testidx<-allidx[-trainidx]  
  
traindata<-mtcarshp[trainidx,]  
testdata<-mtcarshp[testidx,]
```

<https://medium.com/@jayeshbahire/lasso-ridge-and-elastic-net-regularization-4807897cb722>

# OLS Basic Regression

Load data

Split Training

Split Testing

Create Model

Find Training Error

Predict Testdata

Estimate testing error

Run Regularization

Compare

```
lm.model<-lm(hp~mpg+wt+drat,data=traindata)
traindata$train.predicted.hp<-predict(lm.model,
traindata[,c("mpg","wt","drat")])
#training error
train_error<-
rsquare(traindata$hp,traindata$train.predicted.hp)
train_msd<-
sqrt(mean((traindata$hp-traindata$train.predicted.hp)^2))
#test error for lm
testdata$test.predicted.hp<-
predict(lm.model,testdata[,c("mpg","wt","drat")])
test_error<-rsquare(testdata$hp,testdata$test.predicted.hp)
test_msd<-s
qrt(mean((testdata$hp-testdata$test.predicted.hp)^2))
```

```
model_stats<-data.frame(#run_names=c("proportion","rsquare","msd"),
+ lm.train=c(1,train_error,train_msd),
+ lm.test=c(1,test_error,test_msd),stringsAsFactors=F)
```

Model captured 94%/86% of the variation during train/test correspondingly

```
> model_stats
      lm.train  lm.test
1  1.0000000  1.0000000
2  0.9395461  0.8616764
3 41.3058453 48.1315269
```

<https://medium.com/@jayeshbahire/lasso-ridge-and-elastic-net-regularization-4807897cb722>

# OLS Basic Regression

Is it possible to improve?

Load data  
Split Training  
Split Testing  
Create Model  
Find Training Error  
Predict Testdata  
Estimate testing error  
Run Regularization  
Compare

Let us regularization, we need lambda.  
If we know the best lambda, we are done.  
But to find that we create a sequence of lambdas  
And calculate Y for all of them.

```
require(ggplot2)
require(glmnet)
y<-traindata$hp

x<-as.matrix(traindata[,c("mpg","wt","drat")])
lambdas<-10^seq(3,-2,by=-0.1)
glm.fit<-glmnet(x,y,alpha=0,lambda=lambdas)
all_coef<-coef(glm.fit)
betas<-all_coef[2:4,]
fitval<-x%*%betas
```

Fitval is all the y values fitted by the coefficients generated by glmnet.  
We have to find the best fit -- very laborious work

<https://medium.com/@jayeshbahire/lasso-ridge-and-elastic-net-regularization-4807897cb722>

# Cross Validation

Load data  
Split Training  
Split Testing  
Create Model  
Find Training Error  
Predict Testdata  
Estimate testing error  
Run Regularization  
Compare

```
cv.glm.fit<-cv.glmnet(x,y,alpha=0,lambda=lambdas,nfolds=5)
```

```
cv.glm.fit$lambda.min
```

We can start with the lambda and find the best fit.

```
traindata$train.penalized.hp<-predict(  
glm.fit,newx=as.matrix(traindata[,c("mpg","wt","drat")]),  
s=cv.glm.fit$lambda.min,type='response')
```

```
# store results  
model_stats<-cbind(model_stats,  
train.penalized.hp=c(proportion=1,  
rsquare=rsquare(traindata$hp,traindata$train.penalized.hp),  
msd=msd(traindata$hp,traindata$train.penalized.hp)))
```

# OLS – Linear Model

```
> t(model_stats)
```

	proportion	rsquare	msd
lm.train	1.00	0.9317804	43.97826
lm.test	1.00	0.9299146	33.68764
train.penalized.hp	1.00	0.9310128	44.22497
tr.p.hp2	0.98	0.9310381	44.21685
tr.p.hp3	0.88	0.9311583	44.17832
tr.p.hp4	0.80	0.9312467	44.14993
tr.p.hp5	0.78	0.9312689	44.14280
tr.p.hp6	0.76	0.9312911	44.13566
tr.p.hp7	0.66	0.9313948	44.10236
tr.p.hp8	0.50	0.9315417	44.05511
tr.p.hp9	0.40	0.9316201	44.02987
tr.p.hp10	0.20	0.9317353	43.99278
tr.p.hp11	0.10	0.9317680	43.98223
tr.p.hp12	0.05	0.9317771	43.97932
tr.p.hp13	0.00	0.9317803	43.97827
tr.p.hp14	-0.05	0.9317803	43.97827
tr.p.hp15	-0.80	0.9317803	43.97827
tr.p.hp16	-0.90	0.9317803	43.97827
tr.p.hp17	-1.00	0.9317803	43.97827
tr.p.hp18	-2.00	0.9317803	43.97827
tr.p.hp19	-3.00	0.9317803	43.97827



# Run same experiments over test

```
> t(model_stats)
```

	proportion	rsquare	msd
lm.train	1.00	0.9317804	43.97826
lm.test	1.00	0.9299146	33.68764
train.penalized.hp	1.00	0.9310128	44.22497
tr.p.hp2	0.98	0.9310381	44.21685
tr.p.hp3	0.88	0.9311583	44.17832
tr.p.hp4	0.80	0.9312467	44.14993
tr.p.hp5	0.78	0.9312689	44.14280
tr.p.hp6	0.76	0.9312911	44.13566
tr.p.hp7	0.66	0.9313948	44.10236
tr.p.hp8	0.50	0.9315417	44.05511
tr.p.hp9	0.40	0.9316201	44.02987
tr.p.hp10	0.20	0.9317353	43.99278
tr.p.hp11	0.10	0.9317680	43.98223
tr.p.hp12	0.05	0.9317771	43.97932
tr.p.hp13	0.00	0.9317803	43.97827
tr.p.hp14	-0.05	0.9317803	43.97827
tr.p.hp15	-0.80	0.9317803	43.97827
tr.p.hp16	-0.90	0.9317803	43.97827
tr.p.hp17	-1.00	0.9317803	43.97827
tr.p.hp18	-2.00	0.9317803	43.97827
tr.p.hp19	-3.00	0.9317803	43.97827

test.penalized.hp	1.00	0.9435833	30.22462
tst.p.hp2	0.98	0.9433878	30.27694
tst.p.hp3	0.88	0.9423940	30.54154
tst.p.hp4	0.80	0.9415793	30.75673
tst.p.hp5	0.78	0.9413614	30.81406
tst.p.hp6	0.76	0.9411375	30.87281
tst.p.hp7	0.66	0.9399995	31.16983
tst.p.hp8	0.50	0.9380092	31.68259
tst.p.hp9	0.40	0.9366343	32.03200
tst.p.hp10	0.20	0.9335746	32.79624
tst.p.hp11	0.10	0.9318559	33.21781
tst.p.hp12	0.05	0.9309258	33.44374
tst.p.hp13	0.00	0.9300303	33.65982
tst.p.hp14	-0.05	0.9300303	33.65982
tst.p.hp15	-0.80	0.9300303	33.65982
tst.p.hp16	-0.90	0.9300303	33.65982
tst.p.hp17	-1.00	0.9300303	33.65982
tst.p.hp18	-2.00	0.9300303	33.65982
tst.p.hp19	-3.00	0.9300303	33.65982
tst.p.hp20	-4.00	0.9300303	33.65982

WE get to review the R-Squared and MSD



# the test cases...

```
testdata$tst.p.hp18<-predict(glm.fit,newx=as.matrix(testdata[,c("mpg","wt","drat")])),  
s=-2*cv.glm.fit$lambda.min,type='response')  
model_stats<-cbind(model_stats,tst.p.hp18=c(-2.0,rsquare=rsquare(testdata$hp,  
testdata$tst.p.hp18),msd=msd(testdata$hp,testdata$tst.p.hp18)))  
proportion<-c(proportion,-2.0)  
sum((testdata$tst.p.hp18-testdata$hp)^2)
```

```
testdata$tst.p.hp19<-predict(glm.fit,newx=as.matrix(testdata[,c("mpg","wt","drat")])),  
s=-3*cv.glm.fit$lambda.min,type='response')  
model_stats<-cbind(model_stats,tst.p.hp19=c(-3.0,rsquare=rsquare(testdata$hp,  
testdata$tst.p.hp19),msd=msd(testdata$hp,testdata$tst.p.hp19)))  
proportion<-c(proportion,-3.0)  
sum((testdata$tst.p.hp19-testdata$hp)^2)
```

```
testdata$tst.p.hp20<-predict(glm.fit,newx=as.matrix(testdata[,c("mpg","wt","drat")])),  
s=-4*cv.glm.fit$lambda.min,type='response')  
model_stats<-cbind(model_stats,tst.p.hp20=c(-4.0,rsquare=rsquare(testdata$hp,  
testdata$tst.p.hp20),msd=msd(testdata$hp,testdata$tst.p.hp20)))  
proportion<-c(proportion,-4.0)  
sum((testdata$tst.p.hp20-testdata$hp)^2)
```

Regularization and CV are important steps to reduce overfitting

# Other Penalization Methods

- Ridge completed beta-squared all predictors are kept
- Lasso –linear and so some betas are shrunk to zero.
- combination of Ridge and Lasso using elasticnet

# Concluding Remarks

# Let us take stock

*We have run linear regression*

*And we tried to improve using regularization and cross validation. To minimize overfitting....*

*Both CV and Penalization minimize variation not bias.*

*Now we will veer off into logistic regression and develop formalism for various classifiers.*

*RMSE and R-Squared are not applicable.*

*We have to AUC etc to compare Classifier performance.*